

# Training Language Models with Memory Augmentation

Zexuan Zhong<sup>†</sup> Tao Lei\* Danqi Chen<sup>†</sup>

<sup>†</sup>Princeton University

{zzhong, danqic}@cs.princeton.edu, taoleics@gmail.com

## Abstract

Recent work has improved language models remarkably by equipping them with a non-parametric memory component. However, most existing approaches only introduce memories at testing time, or represent them using a separately trained encoder—resulting in sub-optimal training of the language model. In this work, we present TRIME, a novel yet simple training approach designed for training language models with memory augmentation. Our approach uses a training objective that directly takes in-batch examples as accessible memory. We also present new methods for memory construction and data batching, which are used for adapting to different sets of memories—local, long-term, and external memory—at testing time. We evaluate our approach on multiple language modeling and machine translation benchmarks. We find that simply replacing the vanilla language modeling objective by ours greatly reduces the perplexity, without modifying the model architecture or incorporating extra context (e.g., 18.70  $\rightarrow$  17.76 on WikiText-103). We further augment language models with long-range contexts and external knowledge and demonstrate significant gains over previous memory-augmented approaches.<sup>1</sup>

## 1 Introduction

Memory augmentation has become a remarkable approach to enhance language modeling performance without significantly increasing the amount of parameters and computation. By accessing memory units such as a neural cache of recent inputs (Merity et al., 2017; Grave et al., 2017b) and an external look-up table (Khandelwal et al., 2020), a memory-augmented language model (LM) enjoys increased memorization capacity and sets

\*TL currently works at Google Research. The collaboration was initialized before he joined Google.

<sup>1</sup>Our code and pre-trained models are publicly available at <https://github.com/princeton-nlp/TRIME>.

	Training Memory	Testing Memory
vanilla LM	None	None
cont. cache	None	$\mathcal{M}_{\text{local}}$ or $\mathcal{M}_{\text{long}}$
kNN-LM	None	$\mathcal{M}_{\text{ext}}$
TRIMELM	$\mathcal{M}_{\text{local}}$	$\mathcal{M}_{\text{local}}$
TRIMELM <sub>long</sub>	§4.2	$\mathcal{M}_{\text{local}}, \mathcal{M}_{\text{long}}$
TRIMELM <sub>ext</sub>	§4.3	$\mathcal{M}_{\text{local}}, \mathcal{M}_{\text{long}}, \mathcal{M}_{\text{ext}}$

Table 1: A comparison between our TRIME language models and previous approaches: vanilla LM, continuous cache (Grave et al., 2017b,a), kNN-LM (Khandelwal et al., 2020).  $\mathcal{M}_{\text{local}}, \mathcal{M}_{\text{long}}, \mathcal{M}_{\text{ext}}$  denote three types of memories (see §2.2 for more details).

new state-of-the-art records in various language modeling benchmarks.

A major limitation of existing approaches, however, is that the memory units are either introduced at *testing time* (Grave et al., 2017b; Khandelwal et al., 2020) or taken from a *separately* trained model (Yogatama et al., 2021). As a consequence, they are not directly optimized during the training process, resulting in a missed opportunity to achieve even stronger results. In this paper, we pioneer and present a novel yet simple training approach TRIME (Training with In-batch Memories)<sup>2</sup>, that is well-suited for memory augmentation in language modeling. Our approach makes two major departures compared to standard language model training:

**Training objective** Inspired by contrastive representation learning, we propose a training objective that directly leverages in-batch examples as accessible memory. Our training objective is closely connected to neural cache models (Grave et al., 2017b; Merity et al., 2017) and nearest-neighbor language models (Khandelwal et al., 2020), where the next-

<sup>2</sup>We can also interpret TRIME as *three types of memories*, as we will elaborate in the paper.

token probabilities are calculated by comparing encoder outputs against static token embeddings and memory representations. However, previous work only considers incorporating memories at testing time, while we do for both training and testing.

**In-batch memory construction** With this training objective in mind, the key challenge is how to construct memories effectively *during training* while keeping it efficient. We identify three types of memories that can be leveraged at testing time and have been explored in the literature: (a) *local* memory denotes the words that appear in the recent past and are modeled using attention (Vaswani et al., 2017); (b) *long-term* memory<sup>3</sup> denotes long-range context from the same document but cannot be directly accessed due to the limit of input length; (c) *external* memory is used to store the entire training set or any additional corpus (Khandelwal et al., 2020; Borgeaud et al., 2021). To better leverage these memories at testing time, we devise new *data batching* strategies to improve the construction of training memories (§4). By packing consecutive segments from the same document in one training batch, our model can access long-term memories beyond the attention context. Additionally, we pack segments from other documents that have high lexical overlap as a proxy to external memory units. Importantly, these working memories are generated on the fly during training, allowing us to back-propagate to all memory representations.

We instantiate TRIME in three models by considering different sets of training and testing memories (Table 1) and evaluate them on multiple language modeling benchmarks (Merity et al., 2017; Mahoney, 2009), and an IWSLT’14 machine translation task. We highlight our results as follows:

- We first show that we can simply optimize a language model using our training objective *without* long-term and external memory. Without any other modifications, we demonstrate that a 247M Transformer-based model (Vaswani et al., 2017) can achieve improved perplexity from 18.70 to 17.76 on WIKITEXT-103 with negligible overhead. This model can be viewed as a *simple replacement* for vanilla language models.

- By training with consecutive segments in the same batch, our approach is capable of leveraging very *long context* at testing time—up to 15k-25k

<sup>3</sup>Long-term memory may have different interpretations in other contexts and we use *long-term* memory to refer to long-range context in modeling long sequences, following previous work (Martins et al., 2022; Wu et al., 2022).

tokens on WIKITEXT-103 and ENWIK8. Our approach achieves at least competitive performance as previous works (Dai et al., 2019; Martins et al., 2022; Ji et al., 2022) that modify the Transformer architecture to incorporate memories from previous segments, yet our solution is conceptually simpler and computationally cheaper.

- Finally, we train language models by incorporating all other segments in the same batch as memories. Our model works better with a *large datastore* at testing time and improves over the kNN-LM model (Khandelwal et al., 2020) by reducing the test perplexity from 16.23 to 15.47 on WIKITEXT-103. We also demonstrate improvements over the kNN-MT baseline (Khandelwal et al., 2021) on an IWSLT’14 De-En machine translation task.

In summary, we propose a simple approach TRIME for optimizing language models with memory augmentation and demonstrate consistent and significant gains in multiple experimental settings. Our approach does not modify the model architecture and only uses memories at the final prediction step, and hence adds very little computational overhead during both training and inference. As a result, it is compatible with other neural models and techniques such as recurrent networks and compressed attention (Dai et al., 2019; Rae et al., 2020). We hope that our work can encourage the research community to think about better training objectives for language models, given their significant societal impacts (Brown et al., 2020; Chowdhery et al., 2022; Zhang et al., 2022).

## 2 Preliminaries

### 2.1 Language Modeling

In this paper, we mainly focus on improving language models, although we believe that our solutions may extend to most text generation tasks. Later on, we will demonstrate one application in machine translation (§5.4). Neural language models take a sequence of tokens as context  $c_t = x_1, \dots, x_{t-1}$  and map it to a vector representation  $f_\theta(c_t) \in \mathbb{R}^d$ , where  $f_\theta(\cdot)$  is parameterized by a neural network. The next-token probability is:

$$P(w \mid c_t) \propto \exp(E_w^\top f_\theta(c_t)), \quad (1)$$

where  $E_w \in \mathbb{R}^d$  denotes the output embedding of token  $w \in \mathcal{V}$ . The parameters are optimized to minimize the negative log-likelihood of ground truth  $x_t$  during training.

## 2.2 Memory Augmentation

We consider memory as a set of context-target pairs  $\{(c_i, x_i)\}$  following Grave et al. (2017b); Khandelwal et al. (2020). These context-target pairs can be aggregated to obtain the next-token probability weighted by the similarity between hidden representations.<sup>4</sup> We formalize three types of context-target memories as follows:

**Local memory** The local memory is simply the most recent preceding tokens in the same input. Specifically, for  $c_t = x_1, \dots, x_{t-1}$ , it is defined as:

$$\mathcal{M}_{\text{local}}(c_t) = \{(c_j, x_j)\}_{1 \leq j \leq t-1}. \quad (2)$$

Grave et al. (2017b) use the local memory at testing time, denoted by the “continuous cache” model. However, it has been argued less effective for Transformer-based models because they can already learn to leverage recent tokens in the self-attention layers (Khandelwal et al., 2020). Interestingly, we show that using local memory is still beneficial if we consider it during training.

**Long-term memory** Long-term memory denotes long-range context from the same document, but they cannot be directly accessed by attention. For example, if a document contains 10k tokens, only a short segment of text (e.g., 100-1k tokens) can be fed into a Transformer model because the complexity scales quadratically with input length. Formally, we divide a document into consecutive segments  $s^{(1)}, \dots, s^{(T)}$ , where a segment  $s^{(i)}$  contains  $L$  contexts  $s^{(i)} = \{c_1^{(i)}, \dots, c_L^{(i)}\}$ . The long-term memory for  $c_t^{(i)}$  is:

$$\mathcal{M}_{\text{long}}(c_t^{(i)}) = \{(c_j^{(l)}, x_j^{(l)})\}_{1 \leq l < i, 1 \leq j \leq L}. \quad (3)$$

Previous works (Dai et al., 2019; Rae et al., 2020; Martins et al., 2022; Ji et al., 2022; Wu et al., 2022) leverage hidden representations from previous segments with modified Transformer architectures to learn long-range dependency. Our approach does not modify the model architecture and is compatible with these neural architectures. Note that continuous cache can be naturally extended to long-term memory, as we will experiment later.<sup>5</sup>

<sup>4</sup>Other memory-augmented models differ in when the memory was introduced, such as using them in attention, and retrieve texts of different granularity as memory (Guu et al., 2020; Borgeaud et al., 2021).

<sup>5</sup>The earlier continuous cache work was applied to LSTMs on long sequences, as LSTMs can linearly scale with long sequences and there is no need to segment documents.

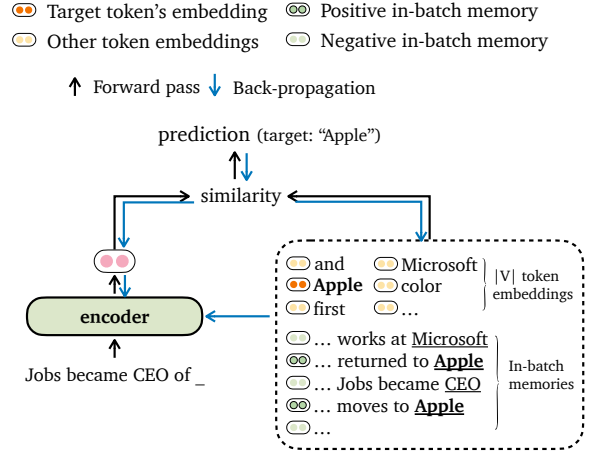


Figure 1: An illustration of our training objective. Our objective aligns the hidden representation with both token embeddings and a set of in-batch contextualized representations that are constructed during training.

**External memory** Finally, external memory assumes a large corpus  $\mathcal{D}$  and the external memory set can be defined as:

$$\mathcal{M}_{\text{ext}} = \{(c_j, x_j) \in \mathcal{D}\}. \quad (4)$$

$\mathcal{D}$  can be simply the training corpus (as is the case in our experiments), or a domain-specific corpus when the testing domain shifts (Khandelwal et al., 2020). Note that  $|\mathcal{M}_{\text{ext}}|$  is usually several orders of magnitude order larger than the previous two types (e.g.,  $10^8$ ); accessing all the memories is computationally expensive and requires approximate nearest neighbor search.

## 3 Training with In-batch Memories

In this section, we propose a new training approach TRIME for language model training. Compared to standard language model training, our training objective assumes a set of *training memories*  $\mathcal{M}_{\text{train}} = \{(c_i, x_i)\}$ . We differentiate training memories from *testing memories*, as they are constructed on-the-fly during training and may deviate from the testing memories used during inference. Importantly, the training memories are constructed from the *same* training batch, which enables back-propagating the training signal to the current hidden representation as well as all the memory representations. We will discuss how to construct training memories in the next section (§4) and only discuss the training objective in a general form.

Our training objective is illustrated in Figure 1. Given a memory set  $\mathcal{M}$  and a context  $c$ , TRIME

defines the next-token probability distribution as:

$$P(w | c) \propto \exp(E_w^\top f_\theta(c)) + \sum_{(c_j, x_j) \in \mathcal{M}_{\text{train}}: x_j = w} \exp(\text{sim}(g_\theta(c), g_\theta(c_j))). \quad (5)$$

Here,  $f_\theta(c)$  is the output representation of a Transformer model and  $E_w$  is the token embedding as we defined in §2.1.  $g_\theta(\cdot)$  denotes the representations that can be used to compute similarity between  $c$  and all the contexts  $c_j$  in the memory  $\mathcal{M}_{\text{train}}$ . It is possible to simply take  $g_\theta = f_\theta$ ; however, we find that taking  $g_\theta$  to be the input of the final feed-forward layer in Transformer works better, which is consistent with the observation in Khandelwal et al. (2020). In addition,  $\text{sim}(\cdot, \cdot)$  is a similarity function and we found using the scaled dot-product  $\text{sim}(q, k) = \frac{q \cdot k}{\sqrt{d}}$  (Vaswani et al., 2017) leads to stable training and better performance in our preliminary experiments.

This training objective can be viewed as a contrastive loss (Hadsell et al., 2006): for a context-target pair  $(c, w^*)$ , the goal is to align the query representation  $f_\theta(c)$  (and  $g_\theta(c)$ ) with the *static* token representation  $E_{w^*}$ , and *contextualized* representations that share the same next token i.e.,  $g_\theta(c_j)$  for  $x_j = w^*$ . Our objective handles rare words nicely—if  $w^*$  does not appear in the training memory, the objective will fall back to aligning  $f_\theta(c)$  with only the word embedding  $E_{w^*}$ . Similar to the vanilla training loss (Eq. 1), our TRIME loss is optimized to minimize the negative log-likelihood of next token  $w^*$  and all the parameters  $\theta$  and  $E_w$  are updated during training.

Our training objective is inspired by the success of contrastive learning in dense retrieval (Karpukhin et al., 2020)—As we will show in §6, it can help improve retrieving contexts that share the same next token effectively when the set of testing memories is large. Our objective is also closely connected to the objective used in Grave et al. (2017b); Khandelwal et al. (2020), which linearly interpolates two distributions: the standard language modeling objective and a distribution defined by cache or an external datastore, e.g.,  $P(w | c) = (1 - \lambda)P_{\text{lm}}(w | c) + \lambda P_{\text{knn}}(w | c)$ . Our work differs from previous works most in that we use this objective as a *training* (and testing) objective, while they only used it *at testing time*—the key is how to construct training memories that we will elaborate next.<sup>6</sup>

<sup>6</sup>Grave et al. (2017b) described a “global normalization”

## 4 Adaption to Different Memories

**Inference** We are interested in incorporating the three types of memories defined in §2.2 and their combinations at testing time. The testing objective is basically the same as the training objective (Eq. 5) except that we take testing memories as a combination of  $\mathcal{M}_{\text{local}}$ ,  $\mathcal{M}_{\text{long}}$  and  $\mathcal{M}_{\text{ext}}$  and we tune a temperature term  $\tau$  to adjust the weight of the memory component. See Appendix A for details about the testing objective.

**Notation** Throughout this section, we use  $L$  to denote segment length,  $B$  to denote the total number of segments used in the one training batch, and  $m$  to denote the number of consecutive segments from each document in the batch. Correspondingly, each batch will contain  $b \approx \frac{B}{m}$  different documents.  $L$ ,  $B$  and  $m$  are hyper-parameters that we will choose for training, and will vary as we consider different memories during inference.

A key challenge is that the testing memories can be very large (e.g.,  $\mathcal{M}_{\text{long}} \sim 10^4$  and  $\mathcal{M}_{\text{ext}} \sim 10^8$  in our experiments) and it is computationally infeasible to keep training memories the same as testing memories. In the following, we will discuss three ways of constructing training memories and data batching, aiming to reduce the discrepancy between training and testing. Along the way, we will also present three major model instantiations: TRIMELM, TRIMELM<sub>long</sub>, TRIMELM<sub>ext</sub> (Table 1), which combine the training strategies and different sets of testing memories.

### 4.1 Local Memory

$\mathcal{M}_{\text{local}}$  only considers all the previous tokens in the same segment. It is straightforward that we can simply use  $\mathcal{M}_{\text{train}} = \mathcal{M}_{\text{local}}$ . As shown in Fig. 2(a), we basically do not need to make *any* modifications compared to standard language model training. All we need is to replace the training objective of Eq. 1 by our objective in Eq. 5, by incorporating  $(c_j, x_j)$ ,  $\forall j < i$  in the memory during both training and testing. The computational overhead is also negligible compared to running neural encoders on the segment  $x_1, \dots, x_L$  itself. We denote this model as TRIMELM, which can be viewed as a lightweight

variant in the paper, which is similar to our objective. However, they only used it at testing time and only considered short-term contexts in calculating the distribution. Another earlier work (Merity et al., 2017) trained a pointer network component with a learned gating component for the interpolation—we attempted training with a similar objective earlier and found it to perform worse than our current objective.



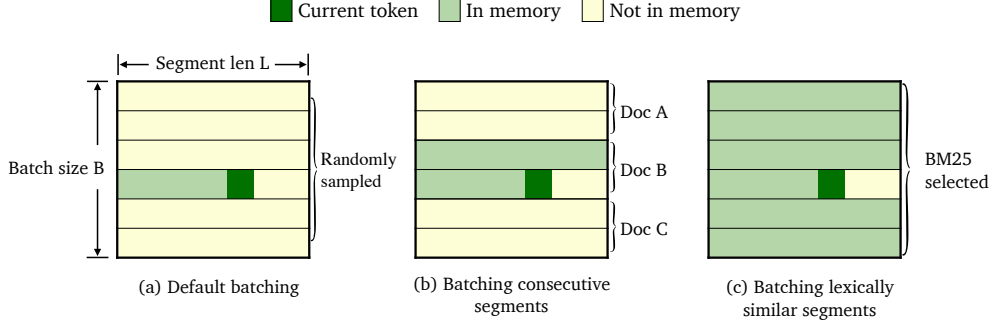


Figure 2: We present several data batching methods and memory construction strategies, in order to adapt to different set of testing memories. (a) default batching: all the segments are randomly drawn from the training corpus (§4.1); (b) we batch consecutive segments from the same document ( $m > 1$ ) in one training batch to better leverage long-range contexts (§4.2); (c) we batch lexically-similar segments in one training batch selected by BM25 to better incorporate a large datastore at testing time (§4.3).

replacement for vanilla language models. As we will show in the experiments, simply incorporating local memory provides a notable gain on multiple LM benchmarks, showing the effectiveness of training with memories explicitly.

## 4.2 Long-term Memory

In order to enable long-term memory augmentation, we pack multiple consecutive segments from the same document in a training batch (i.e.,  $m > 1$ ). For a specific context-target pair  $(c, w)$  in the training batch, its accessible memory  $\mathcal{M}_{\text{train}}$  includes tokens from previous segments as well as the preceding tokens in the same segment. Figure 2(b) illustrates the training batch construction and the training memory for a given token (and its context). Note that at testing time, we can use a much longer context—we simply enumerate the number of segments used in  $\mathcal{M}_{\text{eval}}$  and choose the optimum based on the development set.

We denote this model as  $\text{TRIMELM}_{\text{long}}$ . It shares a similar motivation with many previous works which aim to leverage memory from previous segments through attention recurrence (Dai et al., 2019; Ji et al., 2022), or memory compression (Rae et al., 2020; Martins et al., 2022; Wu et al., 2022). However, our solution deviates significantly from previous approaches. First, previous works need to store the hidden representations (of every layer) from previous segments and modify the self-attention layers to incorporate them. Our approach does not modify the architecture and only uses the outputs from the last layer. In addition, previous works use stale memory representations and do not back-propagate gradients to the representations of previous segments, whereas our batching

method enables gradient propagation to the memory and previous segments.<sup>7</sup> As we will show in the experiments, our approach is competitive with previous works while being conceptually simpler and computationally cheaper.

## 4.3 External Memory

Finally, we consider external memory  $\mathcal{M}_{\text{ext}}$ . Since  $\mathcal{M}_{\text{ext}}$  contains the context-target pairs in a large corpus such as the entire training set, we need to retrieve top- $K$  pairs from  $\mathcal{M}_{\text{ext}}$  measured by  $\text{sim}(g_\theta(c), g_\theta(c_j))$  through (approximate) similarity search (more details are given in §5.2).

Since the retrieved contexts at testing time are expected to be similar to the query context, we propose a simple heuristic for constructing training memories  $\mathcal{M}_{\text{train}}$  by packing segments that have large lexical overlap into the same batch using BM25 (Robertson and Zaragoza, 2009). Specifically, we start with a single segment and repeatedly add segments with highest BM25 scores into the same batch. A high BM25 score indicates that two segments have high lexical overlap and can serve as a good proxy to nearest neighbors in the external memory, which improves our model predictions at testing time. Figure 2(c) illustrates our method.  $\mathcal{M}_{\text{train}}$  contains all tokens from other segments as well as the previous tokens in the same segment. We set  $m = 1$  during training as many segments from the same document tend to have high lexical overlap and denote this model by  $\text{TRIMELM}_{\text{ext}}$ .

In practice, when considering tokens from both the current segment and other segments in the

<sup>7</sup>We also attempted using segments in previous training batches as stale representations and didn’t find any improvement in preliminary experiments.

Model	#Params	Dev ( $\downarrow$ )	Test ( $\downarrow$ )	Speed ( $\uparrow$ )
Transformer (Baevski and Auli, 2019)	247M	17.96	18.65	-
+ continuous cache (Grave et al., 2017b)	247M	17.67	18.27	-
Transformer-XL (Dai et al., 2019)	257M	-	18.30	-
Transformer (our run)	247M	18.04	18.70	3.6k t/s
+ continuous cache	247M	17.65	18.26 (-0.44)	3.6k t/s
TRIMELM	247M	17.10	17.76 (-0.94)	3.6k t/s
TRIMELM <sub>long</sub>	247M	<b>17.01</b>	<b>17.64</b> (-1.06)	3.6k t/s
kNN-LM (our run)	247M	16.40	16.37	300 t/s
+ continuous cache	247M	16.23	16.23 (-0.14)	300 t/s
TRIMELM <sub>ext</sub> (w/o $\mathcal{M}_{\text{long}}$ )	247M	15.62	15.55 (-0.82)	300 t/s
TRIMELM <sub>ext</sub>	247M	<b>15.56</b>	<b>15.47</b> (-0.90)	300 t/s
kNN-LM (Khandelwal et al., 2020) <sup>†</sup>	247M	16.06	16.12	50 t/s
+ continuous cache (Grave et al., 2017b) <sup>†</sup>	247M	15.81	15.79 (-0.33)	50 t/s
TRIMELM <sub>ext</sub> <sup>†</sup>	247M	<b>15.40</b>	<b>15.37</b> (-0.75)	50 t/s

Table 2: Performance of our TRIME models on WIKITEXT-103. <sup>†</sup>: the results are based on computing actual distances instead of using approximated distances returned by FAISS indexes, which requires a large SSD storage. To measure the speed of models (tokens/second), we run the model with a single NVIDIA RTX 3090 GPU and run the FAISS indexer with 32 CPUs.

batch, we observe that the model tends to leverage local memory more and ignore other segments. To encourage the use of information from other segments, we exclude the local memory from  $\mathcal{M}_{\text{train}}$  with a probability of  $p$  during training (we find that  $p = 90\%$  works the best, see §E). This significantly improves performance when the model is evaluated with a large set of external memory.

## 5 Experiments

### 5.1 Datasets and Tasks

We evaluate our approach on two popular language modeling benchmarks: WIKITEXT-103 (Merity et al., 2017) and ENWIK8 (Mahoney, 2009), and a machine translation benchmark: IWSLT’14 DE→EN (see Appendix B for data statistics and detailed task setups; see Appendix C for detailed model configurations).

**WIKITEXT-103** is a word-level language modeling dataset consisting of 103M training tokens. We evaluate on two model configurations—one uses a 247M Transformer model and a segment length  $L = 3,072$  and another one uses a 150M Transformer model with segment length  $L = 150$ .

**ENWIK8** is a character-level language modeling dataset that contains a total of 100M characters. We use a 12-layer Transformer model with a hidden dimension 512 and segment length  $L = 512$ .

**IWSLT’14 DE→EN** is a machine translation task, which consists of 170K translation pairs. We use a Transformer encoder-decoder model. See

Appendix B for how we adjust our approach to the machine translation task.

### 5.2 Training and Inference Details

We implement our approach using the Fairseq library (Ott et al., 2019). For TRIMELM<sub>long</sub> and TRIMELM<sub>ext</sub>, we tune the number of segments used in  $\mathcal{M}_{\text{long}}$  on the development set during evaluation. For our TRIMELM<sub>ext</sub> model which requires building a large datastore at testing time, we use the FAISS library (Johnson et al., 2019) for approximate nearest neighbor search. See our hyperparameters in Appendix C.

We first train our model with the standard LM objective (Eq. 1) for the first 5% updates. Without this warmup stage, we observe the training process to be unstable probably due to a large variance in the estimated distributions. We find that when a large set of external memory  $\mathcal{M}_{\text{ext}}$  is considered during inference, the performance can be improved linearly interpolating the output distribution and a distribution over the memory, similarly to kNN-LM (Khandelwal et al., 2020). Thus, we apply an additional linear interpolation to our output probability distribution when considering external memory  $\mathcal{M}_{\text{ext}}$  (see Appendix F for details).

### 5.3 Results: Language Modeling

We present our language modeling results in Table 2 (WIKITEXT-103, 247M model,  $L = 3,072$ ), Table 3 (WIKITEXT-103, 150M model,  $L = 150$ )

Model	Dev ( $\downarrow$ )	Test ( $\downarrow$ )
Transformer	28.11	29.14
Transformer-XL	23.42	24.56
Compressive Transformer	-	24.41
$\infty$ -former	-	24.22
LaMemo	22.98	23.77
Transformer (our run)	25.31	25.87
+ continuous cache*	22.95	23.59
TRIMELM	24.45	25.60
TRIMELM <sub>long</sub>	<b>21.76</b>	<b>22.66</b>

Table 3: Performance on the WIKITEXT-103 development set (150M model,  $L = 150$ ). TRIMELM<sub>long</sub> uses a long-term memory with 15,000 tokens. Transformer-XL: (Dai et al., 2019), Compressive Transformer: (Rae et al., 2020),  $\infty$ -former: (Martins et al., 2022), LaMemo: (Ji et al., 2022). \*: extending continuous cache to long-term memory.

and Table 4 (ENWIK8).

**TRIMELM vs. vanilla LM** We first compare our TRIMELM model which only uses local memory during training and testing. Table 2 shows that adding a continuous cache during inference can improve the performance of vanilla Transformer from 18.70 to 18.26, and our TRIMELM further improves the perplexity to 17.76. These results suggest that even though the attention mechanism can “see” local context, using local memory during both training and testing can still improve model performance. TRIMELM has no computational overhead compared to vanilla LM (indicated by the “speed” column), making it a simple and better replacement for vanilla language models. Similar trends can be observed in Table 3 and Table 4 (25.87 vs. 25.60 and 1.16 vs. 1.12). The improvement is much smaller though, due to a much smaller segment length  $L$ . More analyses are given in §6.

**TRIMELM<sub>long</sub> leverages long contexts** We then examine our TRIMELM<sub>long</sub> model which is trained with the data batching method described in §4.2. As shown in Table 3 and Table 4, TRIMELM<sub>long</sub> improves vanilla language models substantially (25.87  $\rightarrow$  22.66 on WIKITEXT-103 and 1.16  $\rightarrow$  1.05 on ENWIK8) by leveraging long-range contexts at inference time. We find the model achieves its best results when leveraging 15,000 tokens on WIKITEXT-103 and 24,576 tokens on ENWIK8, even though the segments used during training are much shorter ( $L = 150$  and 512 respectively). We also add continuous cache to the vanilla

Model	#Params	Dev ( $\downarrow$ )	Test ( $\downarrow$ )
T12	44M	-	1.11
Transformer-XL	41M	-	1.06
Adapt-Span	39M	1.04	1.02
Longformer	41M	1.02	1.00
Expire-Span	38M	-	<b>0.99</b>
Transformer (our run)	38M	1.18	1.16
+ continuous cache*	38M	1.16	1.17
TRIMELM	38M	1.14	1.12
TRIMELM <sub>long</sub>	38M	<b>1.08</b>	<b>1.05</b>

Table 4: Performance on the ENWIK8 dataset. TRIMELM<sub>long</sub> achieves the best results by using a long-term memory of a size 24,576. T12: (Al-Rfou et al., 2019), Transformer-XL: (Dai et al., 2019), Adapt-Span: (Sukhbaatar et al., 2019), Longformer: (Beltagy et al., 2020), Expire-Span: (Sukhbaatar et al., 2021). \*: extending continuous cache to long-term memory.

Transformer model and find it to underperform our model, demonstrating the importance of joint training using our approach.

Compared to previous methods which explicitly leverage hidden representations from previous segments (Dai et al., 2019; Rae et al., 2020; Martins et al., 2022; Ji et al., 2022), our approach achieves better or at least competitive performance. Different from these approaches which need to store all the hidden representations of every layer and modify the model architecture, we only incorporate the outputs from the last layer—requiring less computations and GPU memory. We also believe that our approach is orthogonal and can be applied on top of these models and we leave it to future work.

**TRIMELM<sub>ext</sub> vs. kNN-LM** Finally, our model TRIMELM<sub>ext</sub> outperforms the kNN-LM model (Khandelwal et al., 2020), which uses external memory only at testing time—improving the perplexity from 16.23 to 15.47 on WIKITEXT-103 (Table 2). We also evaluate a model which does not use long-term memory (denoted by TRIMELM<sub>ext</sub> w/o  $\mathcal{M}_{\text{long}}$ ) for a fair comparison with kNN-LM with continuous cache and the difference is very small (15.55 vs 15.47). Our results suggest that by using contrastive loss and BM25 batching (§4.3), the model learns to better retrieve and leverage information from a large external memory.

## 5.4 Results: Machine Translation

To showcase the generality of our training approach TRIME to other generation tasks, we evaluate our approach on the IWSLT’14 German-English translation task. Since it is a sentence-level task, we do

Model	BLEU ( $\uparrow$ )
Transformer enc-dec	32.58
kNN-MT	33.15
TRIMELM <sub>ext</sub>	<b>33.40</b>

Table 5: Adapting TRIME to kNN-MT (Khandelwal et al., 2021) on an IWSLT’14 machine translation task.

not use any local or long-term memory ( $\mathcal{M}_{\text{local}}$ ,  $\mathcal{M}_{\text{long}}$ ), as there are few repetitive tokens. We denote our model as TRIMELM<sub>ext</sub>.

As shown in Table 5, our approach improves the vanilla transformer by 0.82 BLEU score and outperforms kNN-MT (Khandelwal et al., 2021). This demonstrates that our approach is able to improve the performance on other language generation tasks with different memory access.

## 6 Analysis

We conduct ablation studies and analyses to further understand individual components of our approach. Due to the limited computation budget, some experiments on WIKITEXT-103 are conducted with a small 7M Transformer model (8 layers, hidden dimension 128) in this section and the trends are generally similar for smaller models (see Appendix C and Appendix D for details).

**Batching and memory construction** We first study how different data batching and memory construction strategies affects the performance when different testing memories are used. We compare our three models (TRIMELM, TRIMELM<sub>long</sub>, TRIMELM<sub>ext</sub>) in Table 6. This ablation study clearly shows that packing consecutive segments and segments with high BM25 scores in the same training batch can improve the performance when the long-range and external memory are used. This demonstrates the importance of closing the gap between training and inference.

**Effectiveness of using local memory** We study the effectiveness of our model TRIMELM that uses only local memory with different segment lengths  $L$ . As shown in table 7, our model significantly outperforms the baselines in all the settings. This suggests that our model can leverage local memory very effectively to improve performance.

**Leveraging long-range contexts** We study if our model is able to handle large long-term memory. As Figure 3 shows, our model is able to effectively handle long-range context (more than 10k

Method	Test memory		
	$\mathcal{M}_{\text{local}}$	$\mathcal{M}_{\text{local}}, \mathcal{M}_{\text{long}}$	$\mathcal{M}_{\text{local}}, \mathcal{M}_{\text{long}}, \mathcal{M}_{\text{ext}}$
TRIMELM	<b>17.10</b>	17.17	17.40
TRIMELM <sub>long</sub>	17.12	<b>17.01</b>	16.48
TRIMELM <sub>ext</sub>	17.69	17.55	<b>15.56</b>

Table 6: Evaluating our three models (w/ different training methods) on different sets of testing memories. The results are based on the development set of WIKITEXT-103 (247M models).

Model	$L = 3072$	$L = 512$	$L = 150$
Transformer	82.70	81.15	81.40
+ cont. cache	67.45	71.29	75.10
TRIMELM	54.89	63.22	71.82

Table 7: Performance on the WIKITEXT-103 development set (7M models). We vary the segment  $L$  here to study the effectiveness of using local memory.

tokens), which goes beyond typical attention context. Compared to continuous cache (Grave et al., 2017b,a), the improvement of our approach becomes larger when more long-term memory is incorporated. This suggests that our model is able to leverage long-range context much more effectively.

### Retrieval performance on external memory

When external memory is used in our experiments, we perform nearest-neighbor search over the entire memory set  $\mathcal{M}_{\text{ext}}$  to retrieve the top  $K$  keys (we use  $K = 1024$ ). Table 9 compares the retrieval accuracy of our approach and kNN-LM (Khandelwal et al., 2020) for different  $K$ . Our approach outperforms kNN-LM in terms of retrieval results; this explains how our final perplexity surpasses kNN-LM when incorporating external memory.

### Perplexity breakdown for different frequencies

Finally, we aim to understand which type of memories improves perplexity of tokens in different frequency groups. We group tokens into 5 buckets according to their frequency on the development set. Table 8 shows the results for different models. Interestingly, TRIMELM and TRIMELM<sub>long</sub> improve the perplexity of rare words (i.e., frequency  $\leq 1k$ ) while achieving similar or slightly worse results for frequent words. TRIMELM<sub>ext</sub> improves perplexity in all the buckets.



Frequency	> 10k	1k-10k	100-1k	10-100	$\leq 10$	avg
Transformer	3.35	4.11	13.63	30.46	240.39	18.04
TRIMELM	3.47	4.15	13.57	28.05	198.33	17.10
TRIMELM <sub>long</sub>	3.43	4.13	13.62	27.89	194.89	17.01
TRIMELM <sub>ext</sub>	<b>3.19</b>	<b>3.85</b>	<b>12.69</b>	<b>25.33</b>	<b>171.74</b>	<b>15.56</b>

Table 8: Averaged perplexity in each frequency bucket on the WIKITEXT-103 development set (247M models).




Figure 3: Performance of TRIMELM<sub>long</sub> (with different segment lengths  $L$ ) and continuous cache (adapted to long-term memory) on the WIKITEXT-103 (7M models) and ENWIK8 development sets.

## 7 Related Work

**Memory-augmented language models** We have discussed continuous cache (Grave et al., 2017b,a), kNN-LM (Khandelwal et al., 2020) and models that leverage representations from long-range context, e.g., (Dai et al., 2019; Rae et al., 2020; Wu et al., 2022) in the previous sections. Yogatama et al. (2021) also aim to combine several types of memories by learning an adaptive gating function; however, their external memory uses a pre-trained vanilla language model. Borgeaud et al. (2021) demonstrate a remarkable performance by augmenting LMs with an external datastore of trillion of tokens and their datastore is built based on chunks of text using off-the-shelf BERT embeddings (Devlin et al., 2019).

**Transformers for long inputs** A large body of research has investigated how to scale self-attention mechanism to long contexts, either through sparse attention (Liu et al., 2018; Child et al., 2019; Beltagy et al., 2020; Zaheer et al., 2020) or sub-quadratic-time attention (Wang et al., 2020; Choro-manski et al., 2020; Peng et al., 2021; Katharopoulos et al., 2020). See Tay et al. (2020) for a com-

Model	Top-1	Top-8	Top-64	Top-1024
kNN-LM	25.82	50.03	69.85	90.84
TRIMELM <sub>ext</sub>	<b>27.64</b>	<b>51.16</b>	<b>70.43</b>	<b>91.19</b>

Table 9: Retrieval performance on external memory of our model and kNN-LM (Khandelwal et al., 2020). We report top- $K$  retrieval accuracy ( $K = 1, 8, 64, 1024$ ).

prehensive survey of efficient Transformers. Our approach is orthogonal, as we only change the training objective and data batching to enable models to use large contexts during inference.

**Retrieval-augmented models for downstream tasks** While our paper focuses on improving language models with memory augmentation, other works improve models on downstream tasks with a retrieval component, such as question answering (de Masson D’Autume et al., 2019; Borgeaud et al., 2021; Guu et al., 2020), dialogue (Fan et al., 2021), and other knowledge-intensive NLP tasks (Lewis et al., 2020; Petroni et al., 2021).

## 8 Conclusion

In this work, we propose TRIME, a training approach for language modeling. Our training objective is inspired by contrastive learning and combines in-batch representations and static token representations. We also present three model instantiations TRIMELM, TRIMELM<sub>long</sub>, TRIMELM<sub>ext</sub>. Through carefully-designed data batching and memory construction during training, we show that our models can leverage long-range contexts and external memory effectively at testing time. Our approach adds little computational overhead and does not modify model architecture, making it compatible with other neural models and techniques. In particular, we believe that our TRIMELM model can be used as an alternative for state-of-the-art auto-regressive language models. For future work, we are interested in training TRIME with large language models and other text generation tasks.

## Acknowledgements

The authors thank Jane Pan, Howard Chen, Alexander Wettig, and the members of Princeton NLP group for helping with proofreading and providing valuable feedback. This research is partially supported by the James M. \*91 Research Innovation Fund for Data Science and a gift from Apple. ZZ is also supported by a JP Morgan PhD fellowship.

## References

- Rami Al-Rfou, Dokook Choe, Noah Constant, Mandy Guo, and Llion Jones. 2019. Character-level language modeling with deeper self-attention. In *Conference on Artificial Intelligence (AAAI)*, volume 33, pages 3159–3166.
- Alexei Baevski and Michael Auli. 2019. Adaptive input representations for neural language modeling. In *International Conference on Learning Representations (ICLR)*.
- Iz Beltagy, Matthew E Peters, and Arman Cohan. 2020. Longformer: The long-document Transformer. *arXiv preprint arXiv:2004.05150*.
- Sebastian Borgeaud, Arthur Mensch, Jordan Hoffmann, Trevor Cai, Eliza Rutherford, Katie Millican, George van den Driessche, Jean-Baptiste Lespiau, Bogdan Damoc, Aidan Clark, et al. 2021. Improving language models by retrieving from trillions of tokens. *arXiv preprint arXiv:2112.04426*.
- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. 2020. Language models are few-shot learners. *Advances in Neural Information Processing Systems (NeurIPS)*, 33:1877–1901.
- Rewon Child, Scott Gray, Alec Radford, and Ilya Sutskever. 2019. Generating long sequences with sparse Transformers. *arXiv preprint arXiv:1904.10509*.
- Krzysztof Choromanski, Valerii Likhoshesterov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, et al. 2020. Rethinking attention with Performers. *arXiv preprint arXiv:2009.14794*.
- Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. 2022. PaLM: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311*.
- Zihang Dai, Zhilin Yang, Yiming Yang, Jaime Carbonell, Quoc V Le, and Ruslan Salakhutdinov. 2019. Transformer-XL: Attentive language models beyond a fixed-length context. *arXiv preprint arXiv:1901.02860*.
- Cyprien de Masson D’Autume, Sebastian Ruder, Lingpeng Kong, and Dani Yogatama. 2019. Episodic memory in lifelong language learning. *Advances in Neural Information Processing Systems (NeurIPS)*, 32.
- Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova. 2019. BERT: Pre-training of deep bidirectional Transformers for language understanding. In *North American Chapter of the Association for Computational Linguistics (NAACL)*.
- Angela Fan, Claire Gardent, Chloé Braud, and Antoine Bordes. 2021. Augmenting Transformers with knn-based composite memory for dialog. *Transactions of the Association of Computational Linguistics (TACL)*, 9:82–99.
- Edouard Grave, Moustapha M Cisse, and Armand Joulin. 2017a. Unbounded cache model for online language modeling with open vocabulary. *Advances in Neural Information Processing Systems (NIPS)*, 30.
- Edouard Grave, Armand Joulin, and Nicolas Usunier. 2017b. Improving neural language models with a continuous cache. In *International Conference on Learning Representations (ICLR)*.
- Kelvin Guu, Kenton Lee, Zora Tung, Panupong Paspapat, and Ming-Wei Chang. 2020. REALM: Retrieval-augmented language model pre-training. In *International Conference on Machine Learning (ICML)*.
- Raia Hadsell, Sumit Chopra, and Yann LeCun. 2006. Dimensionality reduction by learning an invariant mapping. In *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, volume 2, pages 1735–1742.
- Haozhe Ji, Rongsheng Zhang, Zhenyu Yang, Zhipeng Hu, and Minlie Huang. 2022. LaMemo: Language modeling with look-ahead memory. In *North American Chapter of the Association for Computational Linguistics (NAACL)*.
- Jeff Johnson, Matthijs Douze, and Hervé Jégou. 2019. Billion-scale similarity search with GPUs. *IEEE Transactions on Big Data*, 7(3):535–547.
- Vladimir Karpukhin, Barlas Oğuz, Sewon Min, Patrick Lewis, Ledell Wu, Sergey Edunov, Danqi Chen, and Wen-tau Yih. 2020. Dense passage retrieval for open-domain question answering. In *Empirical Methods in Natural Language Processing (EMNLP)*.
- Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. 2020. Transformers are RNNs: Fast autoregressive Transformers with linear attention. In *International Conference on Machine Learning (ICML)*, pages 5156–5165.
- Urvashi Khandelwal, Angela Fan, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. 2021. Nearest neighbor machine translation. In *International Conference on Learning Representations (ICLR)*.

- Urvashi Khandelwal, Omer Levy, Dan Jurafsky, Luke Zettlemoyer, and Mike Lewis. 2020. Generalization through memorization: Nearest neighbor language models. In *International Conference on Learning Representations (ICLR)*.
- Patrick Lewis, Ethan Perez, Aleksandra Piktus, Fabio Petroni, Vladimir Karpukhin, Naman Goyal, Heinrich Küttler, Mike Lewis, Wen-tau Yih, Tim Rocktäschel, et al. 2020. Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems (NeurIPS)*, 33:9459–9474.
- Peter J Liu, Mohammad Saleh, Etienne Pot, Ben Goodrich, Ryan Sepassi, Lukasz Kaiser, and Noam Shazeer. 2018. Generating Wikipedia by summarizing long sequences. In *International Conference on Learning Representations (ICLR)*.
- Matt Mahoney. 2009. Large text compression benchmark.
- Pedro Henrique Martins, Zita Marinho, and André FT Martins. 2022.  $\infty$ -former: Infinite memory Transformer. In *Association for Computational Linguistics (ACL)*.
- Stephen Merity, Caiming Xiong, James Bradbury, and Richard Socher. 2017. Pointer sentinel mixture models. In *International Conference on Learning Representations (ICLR)*.
- Myle Ott, Sergey Edunov, Alexei Baevski, Angela Fan, Sam Gross, Nathan Ng, David Grangier, and Michael Auli. 2019. fairseq: A fast, extensible toolkit for sequence modeling. In *Proceedings of NAACL-HLT 2019: Demonstrations*.
- Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. 2002. BLEU: a method for automatic evaluation of machine translation. In *Association for Computational Linguistics (ACL)*, pages 311–318.
- Hao Peng, Nikolaos Pappas, Dani Yogatama, Roy Schwartz, Noah A Smith, and Lingpeng Kong. 2021. Random feature attention. In *International Conference on Learning Representations (ICLR)*.
- Fabio Petroni, Aleksandra Piktus, Angela Fan, Patrick Lewis, Majid Yazdani, Nicola De Cao, James Thorne, Yacine Jernite, Vladimir Karpukhin, Jean Maillard, et al. 2021. KILT: a benchmark for knowledge intensive language tasks. In *North American Chapter of the Association for Computational Linguistics (NAACL)*, pages 2523–2544.
- Matt Post. 2018. A call for clarity in reporting BLEU scores. In *Proceedings of the Third Conference on Machine Translation: Research Papers*, pages 186–191, Belgium, Brussels. Association for Computational Linguistics.
- Jack W Rae, Anna Potapenko, Siddhant M Jayakumar, and Timothy P Lillicrap. 2020. Compressive Transformers for long-range sequence modelling. In *International Conference on Learning Representations (ICLR)*.
- Stephen Robertson and Hugo Zaragoza. 2009. *The probabilistic relevance framework: BM25 and beyond*. Now Publishers Inc.
- Sainbayar Sukhbaatar, Édouard Grave, Piotr Bojanowski, and Armand Joulin. 2019. Adaptive attention span in Transformers. In *Association for Computational Linguistics (ACL)*, pages 331–335.
- Sainbayar Sukhbaatar, Da Ju, Spencer Poff, Stephen Roller, Arthur Szlam, Jason Weston, and Angela Fan. 2021. Not all memories are created equal: Learning to forget by expiring. In *International Conference on Machine Learning (ICML)*, pages 9902–9912. PMLR.
- Yi Tay, Mostafa Dehghani, Dara Bahri, and Donald Metzler. 2020. Efficient Transformers: A survey. *arXiv preprint arXiv:2009.06732*.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. Attention is all you need. *Advances in Neural Information Processing Systems (NIPS)*, 30.
- Sinong Wang, Belinda Z Li, Madian Khabza, Han Fang, and Hao Ma. 2020. Linformer: Self-attention with linear complexity. *arXiv preprint arXiv:2006.04768*.
- Yuhuai Wu, Markus N Rabe, DeLesley Hutchins, and Christian Szegedy. 2022. Memorizing Transformers. In *International Conference on Learning Representations (ICLR)*.
- Dani Yogatama, Cyprien de Masson d’Autume, and Lingpeng Kong. 2021. Adaptive semiparametric language models. *Transactions of the Association of Computational Linguistics (TACL)*, 9:362–373.
- Manzil Zaheer, Guru Guruganesh, Kumar Avinava Dubey, Joshua Ainslie, Chris Alberti, Santiago Ontanon, Philip Pham, Anirudh Ravula, Qifan Wang, Li Yang, et al. 2020. Big Bird: Transformers for longer sequences. In *Advances in Neural Information Processing Systems (NeurIPS)*, volume 33, pages 17283–17297.
- Susan Zhang, Stephen Roller, Naman Goyal, Mikel Artetxe, Moya Chen, Shuohui Chen, Christopher Dewan, Mona Diab, Xian Li, Xi Victoria Lin, et al. 2022. OPT: Open pre-trained Transformer language models. *arXiv preprint arXiv:2205.01068*.
- Xin Zheng, Zhirui Zhang, Junliang Guo, Shujian Huang, Boxing Chen, Weihua Luo, and Jiajun Chen. 2021. Adaptive nearest neighbor machine translation. In *Association for Computational Linguistics (ACL)*, pages 368–374.

## A Inference Method

Formally speaking, our testing objective is basically the same as the training objective (Eq. 5):

$$P(w | c) \propto \exp(E_w^\top f_\theta(c)) + \sum_{(c_j, x_j) \in \mathcal{M}_{\text{eval}}: x_j = w} \exp\left(\frac{\text{sim}(g_\theta(c), g_\theta(c_j))}{\tau}\right), \quad (6)$$

except that we take  $\mathcal{M}_{\text{eval}}$  as a combination of  $\mathcal{M}_{\text{local}}$ ,  $\mathcal{M}_{\text{long}}$  and  $\mathcal{M}_{\text{ext}}$ . Because  $\mathcal{M}_{\text{eval}}$  may be different from the training memories, we tune a temperature term  $\tau$  to adjust the weight of the memory component when calibrating the distribution, based on the development set.

## B Dataset Statistics and Tasks

We evaluate our approach on three benchmarks: WIKITEXT-103, ENWIK8, and IWSLT’14. Table 10 shows the statistics.

**WIKITEXT-103** is a word-level language modeling dataset consisting of 103M training tokens. Following standard practice, we use adaptive softmax and adaptive token embeddings (Baeviski and Auli, 2019) in our model and report perplexity. In order to better compare with previous work, we evaluate on two model configurations—one uses a 247M Transformer model and a segment length  $L = 3,072$  following Baeviski and Auli (2019); Khandelwal et al. (2020) and another one uses a 150M Transformer model with segment length  $L = 150$  following Dai et al. (2019). More details are provided in Appendix C.

**ENWIK8** is a character-level language modeling dataset that contains a total of 100M characters. Following previous work, we report bit-per-character (bpc) on this dataset. We use a 12-layer Transformer model with a hidden dimension 512 and segment length  $L = 512$ .

We also evaluate the **IWSLT’14** DE→EN machine translation task, which consists of 170K translation pairs. Following Khandelwal et al. (2021), we build an external memory by taking all the translation contexts and the corresponding target token  $((x, y_{<t}), y_t)$  on the training set. We use the output representation as  $f((x, y_{<t}))$  and the input representation of last FFN layer as  $g((x, y_{<t}))$  to compute the loss. Similarly, we use BM25 to batch training data – we encourage two target sentences with a high BM25 score to be in the same training batch. We use the default model configuration

in the Fairseq library (Ott et al., 2019), and sacrebleu (Post, 2018) to compute BLEU scores (Papineni et al., 2002).

	Train	Dev	Test	$ \mathcal{V} $	len
WIKITEXT-103	110M	0.2M	0.3M	270K	3.6K
ENWIK8	94M	5.2M	5.2M	256	-
IWSLT’14	160K	7K	6K	16K	-

Table 10: Statistics of the 3 datasets used in our paper. WIKITEXT-103 is a word-level LM task and ENWIK8 is a character-level language modeling task, and IWSLT’14 is a German-English machine translation task. len: denotes the average document (sentence) length in each dataset. IWSLT’14 is a sentence-level task, so incorporating long-range context will not help.

## C Model Configurations and Hyperparameters

Table 11 shows the model configurations and hyperparameters that we used in our experiments. Following Baeviski and Auli (2019), during training, we train the model with fixed-length segments; during evaluation, we evaluate on the tokens at the end of the segment (i.e., an evaluation segment can overlap with others).

When evaluating with large external memory, we always retrieve  $K = 1,024$  context-target pairs for language modeling. For machine translation, we tune  $K = \{1, 2, 4, 8, 16, 32, 64\}$  following Zheng et al. (2021).

## D Performance of the 7M model on WIKITEXT-103

We conduct our ablation studies and analyses in §6 with an 8-layer Transformer model due to the limited computation budget. The model consists of 8 layers and 4 heads in each layer. The embedding dimension is 128 and the intermediate dimension is 512. The model takes a segment of 3072 tokens as input. We compare our approach with baselines on this model architecture. As shown in Table 12, our approach improves over the baselines by a large margin. This shows that modeling memory explicitly is essential when the model capacity is limited.

## E Tuning $p$ for Training with External Memory

When training model with local and external memory, to avoid the model to only relies on high-quality local memory, we disable the local memory



Dataset	WIKITEXT-103			ENWIK8	IWSLT'14
<b>Model</b>					
#Params	247M	150M	7M	38M	39M
#Layers	16	16	8	12	6+6
Hidden dimension	1024	410	128	512	512
FFN intermediate dimension	4096	2100	512	2048	1024
Adaptive softmax?	yes	yes	yes	no	no
<b>Training</b>					
Segment length	3072	150	3072	512	-
#Tokens per update	73728	36000	24576	49152	16384
Gradient accumulation	6	1	1	1	1
Batch size per update	24	240	8	96	-
#Consecutive segments	4	240	8	96	-
#Long-term memories	12288	36000	24576	49152	-
<b>Evaluation</b>					
Segment length	512	64	512	80	-
#Optimal long-term memories	12288	15000	12288	24576	-
<b>Optimizer and scheduler</b>					
Optimizer type	nag	adam	adam	adam	adam
Learning rate	1.0	2.5e-4	5e-4	2.5e-4	5e-4
Grad crop norm	0.1	0.0	0.0	0.0	0.0
Update steps	286000	400000	200000	400000	170000
Scheduler type	cosine	cosine	inverse_sqrt	cosine	cosine
Linear warmup steps	16000	0	8000	0	4000

Table 11: Model configurations and hyperparameters in our experiments.

Model	Dev	Test
Transformer	82.68	83.66
+ Continuous Cache	67.45	68.08
kNN-LM	51.88	52.24
+ Continuous Cache	45.15	45.82
TRIMELM	54.78	54.69
TRIMELM <sub>long</sub>	60.71	60.10
TRIMELM <sub>ext</sub>	<b>38.51</b>	<b>39.90</b>

Table 12: Performance of the 7M Transformer model on the WIKITEXT-103 dataset.

with a probability of  $p$ . Here we study how  $p$  will affect the final performance of our model. The results of using different  $p$  are shown in Table 13. We find that when  $p = 0$ , the model performs poorly with external memory as the model learns to only leverage local memory and ignores external memory during training. By increasing  $p$ , this issue is mitigated. We set  $p = 0.9$  in our main experiments.

$p$	0.0	0.1	0.5	0.9	1.0
Perplexity	54.33	49.85	44.10	41.50	41.86

Table 13: The performance of our model TRIMELM<sub>ext</sub> on the development set (WIKITEXT-103, 7M model). We disable the local memory with a probability of  $p$  during training.

## F Linear Interpolation When Using $\mathcal{M}_{\text{ext}}$

We find that when a large set of external memory  $\mathcal{M}_{\text{ext}}$  is considered during inference, the performance can be improved by calibrating a separated distribution over the memory and interpolating the output distribution and the memory distribution, similarly to kNN-LM (Khandelwal et al., 2020). We think this is because the distribution of the similarity values has been significantly shifted during inference, while the relative ranking preserves. As a result, having values from two different distribution in one softmax normalization is sub-optimal compared to computing two separated probabilities and interpolating them. We think this is because the

distribution of the similarity values has been significantly shifted during inference, while the relative ranking preserves.

Thus, we apply an additional linear interpolation to our output probability distribution. Specifically, we first use Eq. 6 to compute the distribution  $P(w \mid c)$ . Then, we compute a probability distribution over the tokens in memory  $P'(w \mid c)$  as follow,

$$P'(w \mid c) \propto \sum_{(c_j, x_j) \in \mathcal{M}_{\text{eval}}: x_j = w} \exp\left(\frac{\text{sim}(g_\theta(c), g_\theta(c_j))}{\tau'}\right). \quad (7)$$

We linearly interpolate these two probability distributions with a coefficient  $\lambda$  and get the final output  $P_{\text{final}}(w \mid c)$ :

$$P_{\text{final}}(w \mid c) = (1 - \lambda)P(w \mid c) + \lambda P'(w \mid c). \quad (8)$$

We tune the temperature terms and  $\lambda$  on the development set.